


Lecture 8

Google's Map/Reduce

EN 600.320/420

Instructor: Randal Burns

5 March 2014



Department of Computer Science, *Johns Hopkins University*

Overview

- Map/Reduce
 - A structured programming model for data parallelism
 - Geared toward data intensive applications
- The Google File System
 - Large scale parallel file system for low \$\$
 - Failures are the normal case
- Open-source re-implementation
 - HADOOP: Map/Reduce
 - HDFS: Google File System



Systems that Use Map/Reduce

- Compile to Apache Hadoop!
 - PIG, Hive
- That build on HDFS
 - Hbase
- Systems that run Hadoop! on their storage
 - Cassandra, MongoDB, Aster, Oracle NoSQL



Map/Reduce Model

- Map
 - Process key/value pairs
 - Produce intermediate key/value pairs
- Reduce
 - Merge intermediate pairs on key value
- Map and reduce were LISP primitives
 - Hence the functional style
 - Actually write procedural code under functional constraints



Example: Count all Words

- This is pseudocode
 - Need a tokenizer, etc.

```
map(String key, String value):  
    // key: document name  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");
```



Example: Count all Words

```
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```



What can you do?

- Seems like a limited model
- But, many string processing problems fit naturally
- Can be used iteratively
- Examples
 - Grep
 - Web-log processing (e.g. URL hit counts)
 - Reverse Web-link graph
 - Terms vector per host
 - Build an inverted index
 - Distributed sort (naturally)



Types and Domains

- Map is a transformation
 - Input domain to output domain
- Reduce is a collection
 - No domain change
- C++ implementation is based all on strings
 - User code must convert to structured types

```
map      (k1, v1)      → list (k2, v2)
reduce   (k2, list (v2)) → list (v2)
```



The Why of Map/Reduce

- *How does programming in Map/Reduce help parallelism?*
 - Consider the count all words pseudo-code example



The Why of Map/Reduce

- *How does programming in Map/Reduce help parallelism?*
 - Consider the count all words pseudo-code example
- Potential parallelism
 - Map: on all the sites that store data
 - Or on all the sites to which you distribute data after accessing it from storage
 - Reduce: as many computers as there are terms

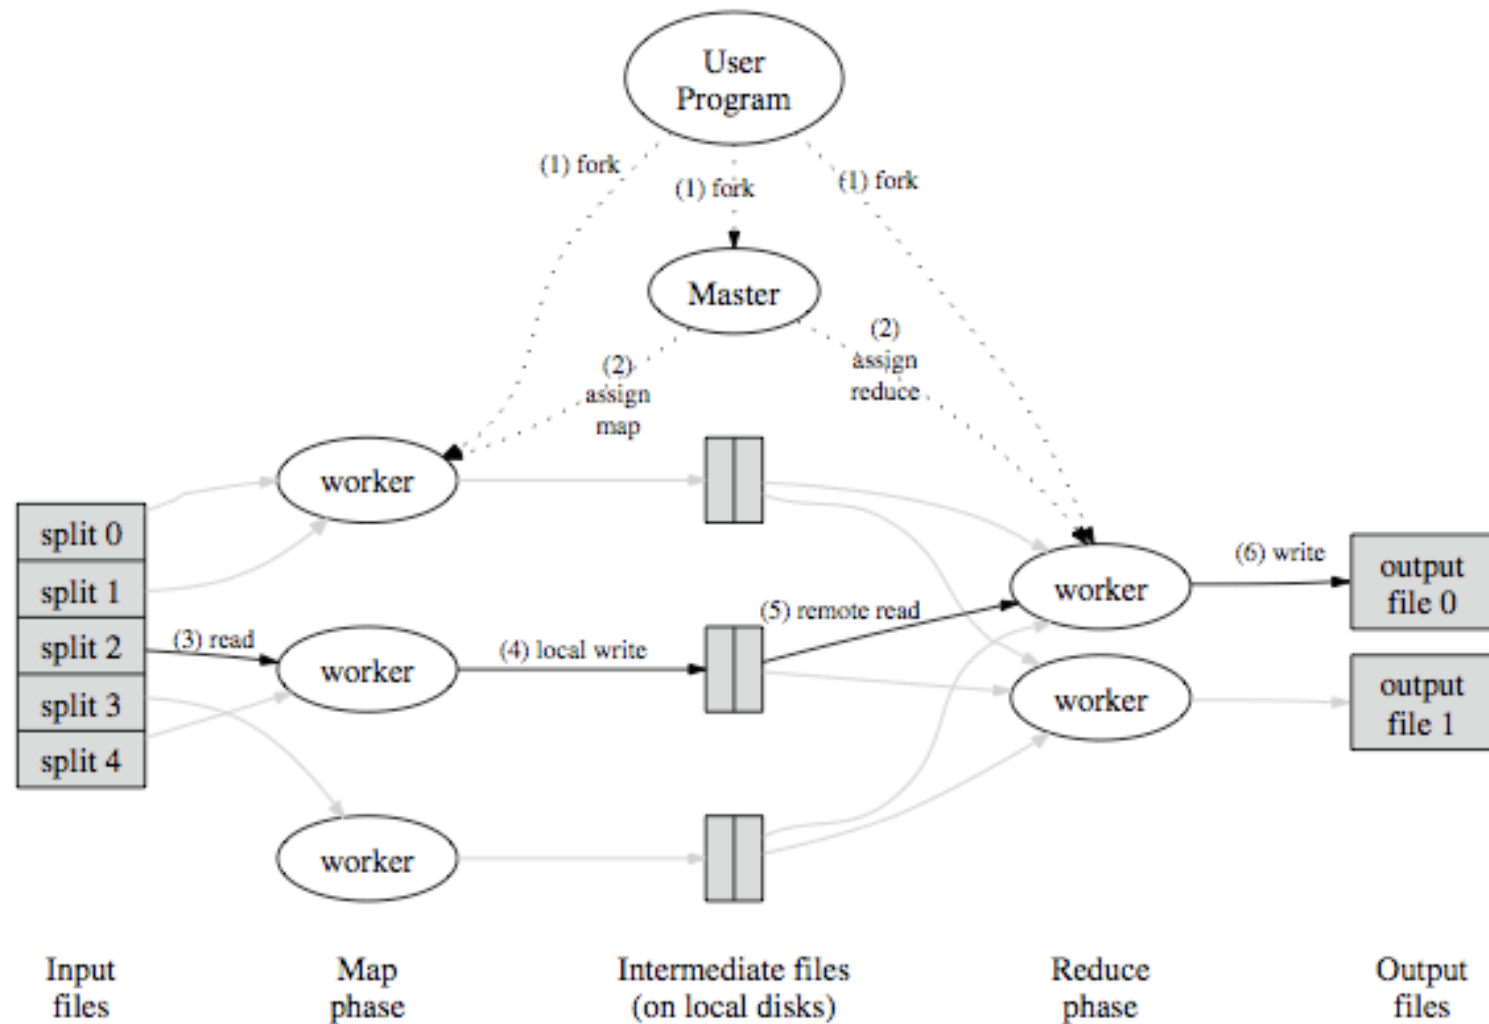


The Point

- Programs in Map/Reduce are automatically parallelized by the Map/Reduce runtime system
- No programmer interaction with parallelism
 - Except encoding problem well in Map/Reduce
- No explicit knowledge of
 - # of machines
 - Location of machines
- Scale up parallelism for arbitrary configurations
- (For embarrassingly parallel problems?)



Map/Reduce Runtime



Runtime Properties

- Automatically partition input data
 - 16-64 MB chunks for Google
- Create M map tasks: one for each chunk
 - Assign available workers (up to M) to tasks
- Write intermediate pairs to local (to worker) disk
- R reduce tasks (defined by user) read and process intermediate results
- Output is R files available on shared file system
- Master tracks state
 - Assignment of M map tasks and R reduce tasks to workers
 - State and liveness of the above



System Issues

- What kind of problems arise with which the runtime system needs to deal?



System Issues

- Master failure
 - Checkpoint/restart, classic distributed systems/replication problem
- Failed worker
 - Heartbeat liveness detection, restart
- Slow worker
 - Backup tasks
- Locality of processing to data
 - Big deal, they don't really solve
- Task granularity
 - Metadata size and protocol scaling (not inherent parallelism) limit the size of M and R



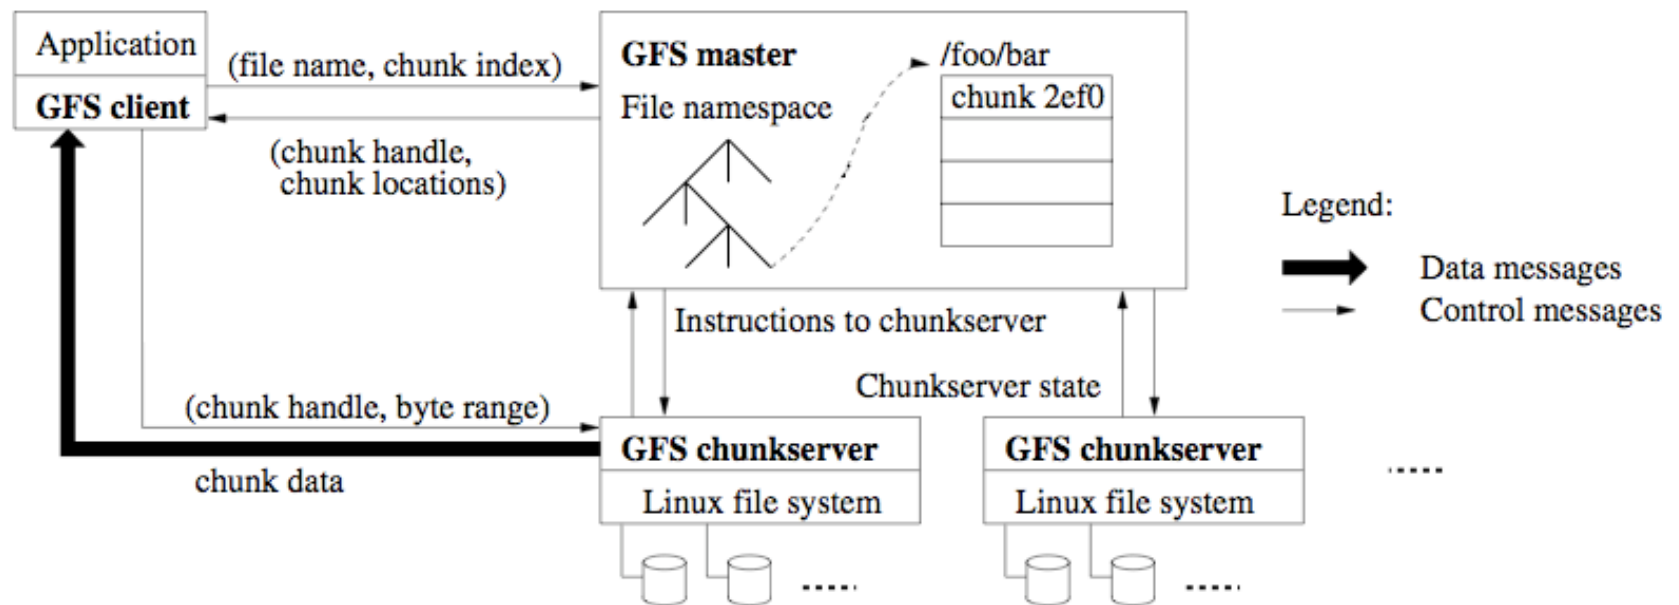
Google File System

- GFS
- Same goals
 - Wide-distribution
 - Commodity hardware
 - High (aggregate) performance
- Different environment?
 - Component failures are normal behavior
 - Files are huge (specific to Google environment)
 - Most files have append-only writing



Architecture

- Looks like every other out-of-band FS



What's Cool

- Atomic checkpoint and append
 - Major mode for writing
 - Great semantics for limited usage
- In-memory metadata at Master
 - Gotta keep it small
- Re-replication
 - Keep three, detect missing on read/write

