


Lecture 12

Safety and Liveness

EN 600.320/420

Instructor: Randal Burns

31 March 2014



Department of Computer Science, *Johns Hopkins University*

Characterizing Concurrency

- Safety (correctness): what are the semantic guarantees (invariants) expressed by a locking protocol under concurrent execution
 - Typically related to some notion of serial execution
- Liveness (progress): how can the execution of one thread be delayed by other threads
 - **Starvation Freedom:** *If a process is trying to enter its critical section, then this process must eventually enter its critical section*
 - Many algorithms ignore starvation freedom on the premise that contention is rare



Ideal Properties

Safety

- FIFO: all operations execute serially
 - Concurrent execution does not change the semantics of computing

Liveness

- Non-blocking




Waiting

- *Doorway* is *wait free*
 - Bounded number of atomic steps
- Waiting is *busy waiting* on some condition
- Notions of waiting:

r-bounded: a process will enter its critical section before every other process executes $r+1$ critical sections

- **FIFO** is 0-bounded waiting



The Bakery Algorithm

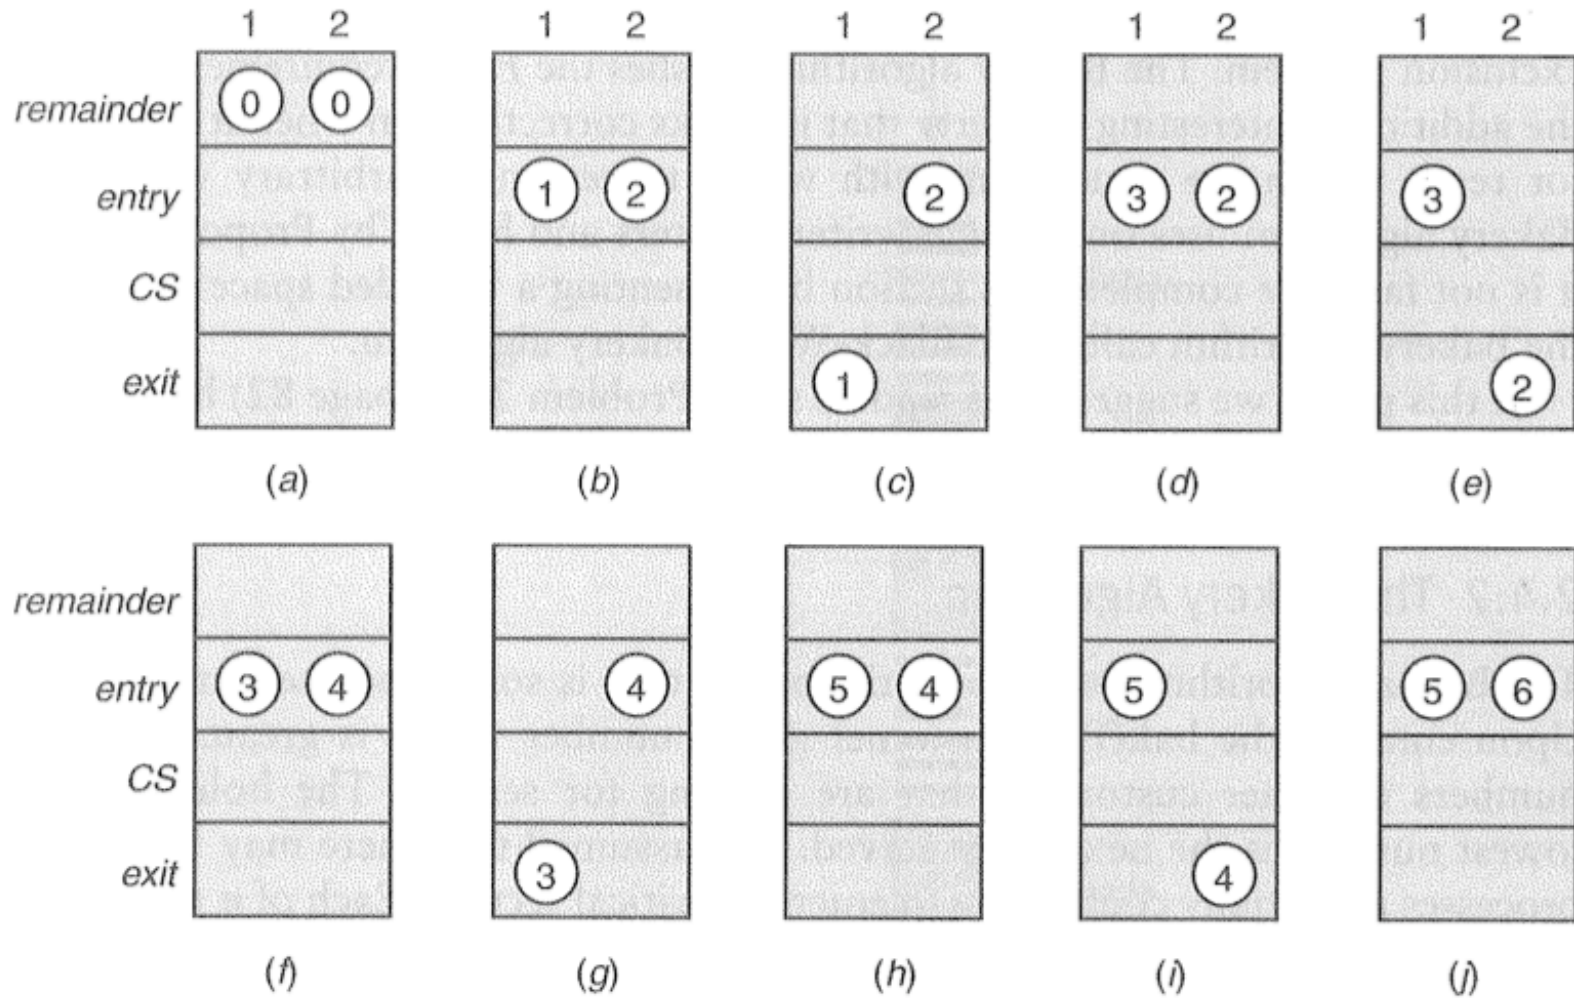
- FIFO processing for Mutual Exclusion

Initially: all entries in *choosing* and *number* are *false* and 0, respectively.

```
1  choosing [i] := true;  
2  number [i] := 1 + maximum(number [1], ..., number [n]);  
3  choosing [i] := false;  
4  for j = 1 to n do  
5      await choosing [j] = false;  
6      await (number [j] = 0 or (number [j], j) ≥ (number [i], i))  
7  od;  
8  critical section;  
9  number [i] := 0;
```



The Bakery Algorithm



Observations about Bakery

- Code parts
 - Lines 1-3 are doorway
 - Lines 4-7 are waiting
 - Line 9 is exit
- Boolean array *choosing[j]* and integer array *number[j]*
 - Read by all processes
 - Written only by process j
- Uses lexicographic ordering of nodes as well as ticket numbers
 - $(a,b) < (c,d)$: $a < c$ or $(a = c \text{ and } b < d)$



Properties of Bakery

- FIFO: no process that enters the *doorway* gets ahead of a process that has already started *waiting*
- Satisfies mutual exclusion
- Is not fast!!!
 - $3(n-1)$ memory accesses when there is no contention
 - This is why Fast Mutual Exclusion is so cool
 - Exchange FIFO for constant overhead
 - Don't read other processes state if there is not contention



The Next Layer of Concepts

- Variants on number of processes
 - Infinitely many processes
 - Sparse process id address space (symmetric algs.)
- Spinning on local registers only
- For different memory models
 - CC, DSM



Building on Primitives

- Common atomic operations (building blocks):
 - Read
 - Write
 - Test-and-set
 - Swap
 - Fetch and add (fetch and increment)
 - Read-modify-write
 - Compare-and-swap



Test-and-Set Bit

- Two operations
 - Reset: write 0
 - Test and set: write 1 and return old value
- Trivial deadlock free synchronization

```
await (test-and-set(x) = 0);  
critical section  
reset(x);
```

- This is called a *spin lock*
 - Mutual exclusion, deadlock free
 - Not starvation resistant



Test-and-Test-and-Set Bit

- Test-and-set alg. writes bit every iteration
 - Invalidates caches even when data don't change
- Test-and-test-and-set
 - Supports test w/out set
- Produces fewer cache misses
 - What's the miss pattern during contention

```
await (x=0);  
while (test-and-set(x) = 1) do  
    await (x=0) od;  
critical section  
reset(x);
```



What's wrong with Spin Locks?

- Every process spins on shared state
- When lock is freed, all processes attempt to acquire
- Performance varies with contention:
 - Low contention good (simple algorithms)
 - High contention bad (burst of activity: messages and cache invalidations)
- Can be addressed with backoff policies
 - Like exponential backoff in TPC
- But, queuing is better



Ticket Algorithm

- Bakery algorithms using read-modify-write
 - $<$ and $>$ indicate RMW boundaries

THE TICKET ALGORITHM: process i 's program.

```
constant  $N = \{0, 1, \dots, n - 1\}$ 
shared  $(ticket, valid)$  a read-modify-write register ranges over  $N \times N$ ;
initially  $ticket = valid$ 
local  $(ticket_i, valid_i)$  ranges over  $N \times N$ 

1   $\langle (ticket_i, valid_i) := (ticket, valid) \rangle$ ;
2   $ticket := (ticket + 1) \bmod n$ ;
3  while  $ticket_i \neq valid_i$  do
4       $\langle valid_i := valid \rangle$  od;
5  critical section;
6   $\langle valid := (valid + 1) \bmod n \rangle$ ;
```



Properties of RMW Ticket Alg.

- FIFO: in the order of successful RMW
- Mutual exclusion and deadlock freedom
- Uses one shared register that holds n^2 values
- This is the power of H/W support
 - Modern processors provide some variant of RMW




Waiting w/out the Busy Wait

- The Semaphore S
 - $\text{up}(S)$ increase the value of S
 - $\text{down}(S)$ decrease the value of S
 - Binary semaphore takes values 0 and 1
- Using the semaphore
 - $\text{down}(S)$; critical section; $\text{up}(S)$;
 - To realize deadlock-free, mutual exclusion
- Where does the busy wait go?
 - Nowhere: implement semaphores with test-and-set
 - Into the kernel: one process does all the busy waiting
 - Into hardware: use interrupts



Barriers

- Allows a “synchronous” algorithm to run on asynchronous hardware



Simple Barrier

- Built on an atomic counter and atomic bits

```
shared  counter: atomic counter ranges over  $\{0, \dots, n\}$ , initially 0
        go: atomic bit, initial value is immaterial
local   local.go: a bit, initial value is immaterial

1 local.go := go                                /* remembers current value */
2 counter := counter + 1
        /* atomically increment the counter */
3 if counter = n then                          /* last to arrive to the barrier */
4   counter := 0                               /* reset the barrier */
5   go := 1 - go                               /* notify all */
6 else await(local.go  $\neq$  go) fi                /* not the last to arrive */
```



Multiple Resources

- To now, we have talked about deadlock freedom for mutual exclusive access to a single resource
- With multiple resources, we get deadlock even with deadlock-free access to each resource



Deadlock

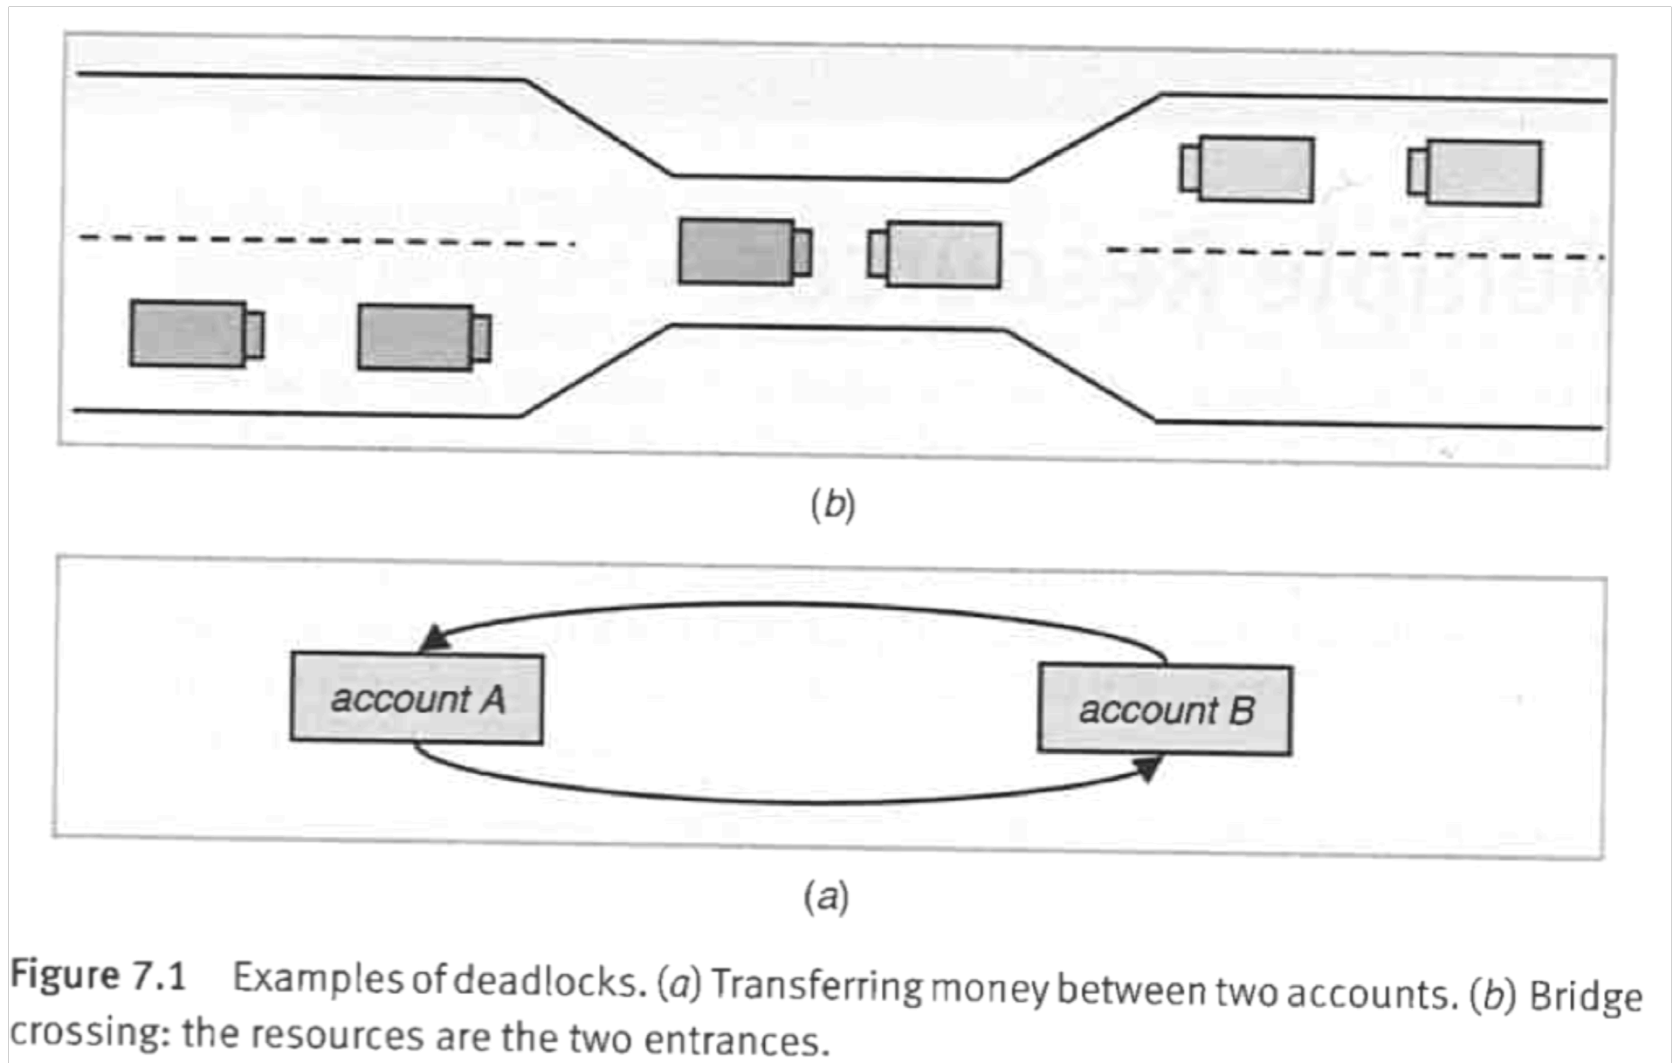


Figure 7.1 Examples of deadlocks. (a) Transferring money between two accounts. (b) Bridge crossing: the resources are the two entrances.



Deadlock

- *Def'n:* A set of processes are deadlocked if each process in the set is waiting for an event that only another process in the set can cause
- Requirements (all four must hold simultaneously)
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait
- We'll do more on deadlock in MPI

